

February, 2024

---



# Health Insurance premium predictive analysis

Course: Quantitative Methods for  
Analytics

Instructor: Leslie Major

---

Prepared by  
**Battsetseg Tolya**  
Student ID: A01407044

## Table of contents

1. Introduction.....	page 1
2. Understanding the data and data cleaning.....	page 2-3
3. Exploratory data analysis.....	page 3-6
4. Comparison of predictive models.....	page 7-10
5. Executive summary .....	page 11
6. Appendix.....	page 12

## Introduction

Based on given hospital treatment charges of patients and other related information about patients, such as age, gender, BMI, region, number of dependents, and whether the patient is a smoker or not, health insurance companies can predict possible hospital charges of the patient based on factors mentioned above, so it can change its premium fee based on these factors and customize their premium fee to decrease their expenses. If predicted hospital charges tend to be higher, then health insurance companies can charge a higher premium fee when having a contract with the insured. When health insurance companies predict hospital charges and customize premium fees, it can save money and help the insurance company to work with a profit.

## Understanding of data and data cleaning

The data has been taken from [www.kaggle.com](http://www.kaggle.com) and the data has a following independent variables:

- age: age of primary beneficiary
- sex: insurance contractor gender: female or male
- bmi: Body mass index, providing an understanding of body, weights that are relatively high or low relative to height, objective index of body weight ( $\text{kg} / \text{m}^2$ ) using the ratio of height to weight, ideally 18.5 to 24.9
- children: Number of children covered by health insurance / Number of dependents
- smoker: Smoking
- region: the beneficiary's residential area in the US, northeast, southeast, southwest, northwest.
- charges: Individual medical costs billed by health insurance, which is the target variable.

```
print(data)
```

	age	gender	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

```
[1338 rows x 7 columns]
```

In Python, I did the following commands to see my data:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   gender      1338 non-null   object
 2   bmi         1333 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
data.describe()
```

```
[59]:
```

	age	bmi	children	charges
count	1338.000000	1333.000000	1338.000000	1338.000000
mean	39.207025	30.658545	1.094918	13270.422265
std	14.049960	6.092785	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.315000	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.675000	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

Based on the tables above, there are five BMI records has missing values. Handled the missing data by putting its averages in the place of empty cells.

```
[60]:
```

```
#handling missing values
count_nan = data.isnull().sum() # the number of missing values for every column
print(count_nan[count_nan > 0])
```

```
bmi    5
dtype: int64
```

```
[61]:
```

```
#fill in the missing values
data.fillna({'bmi':data['bmi'].mean()}, inplace=True)
```

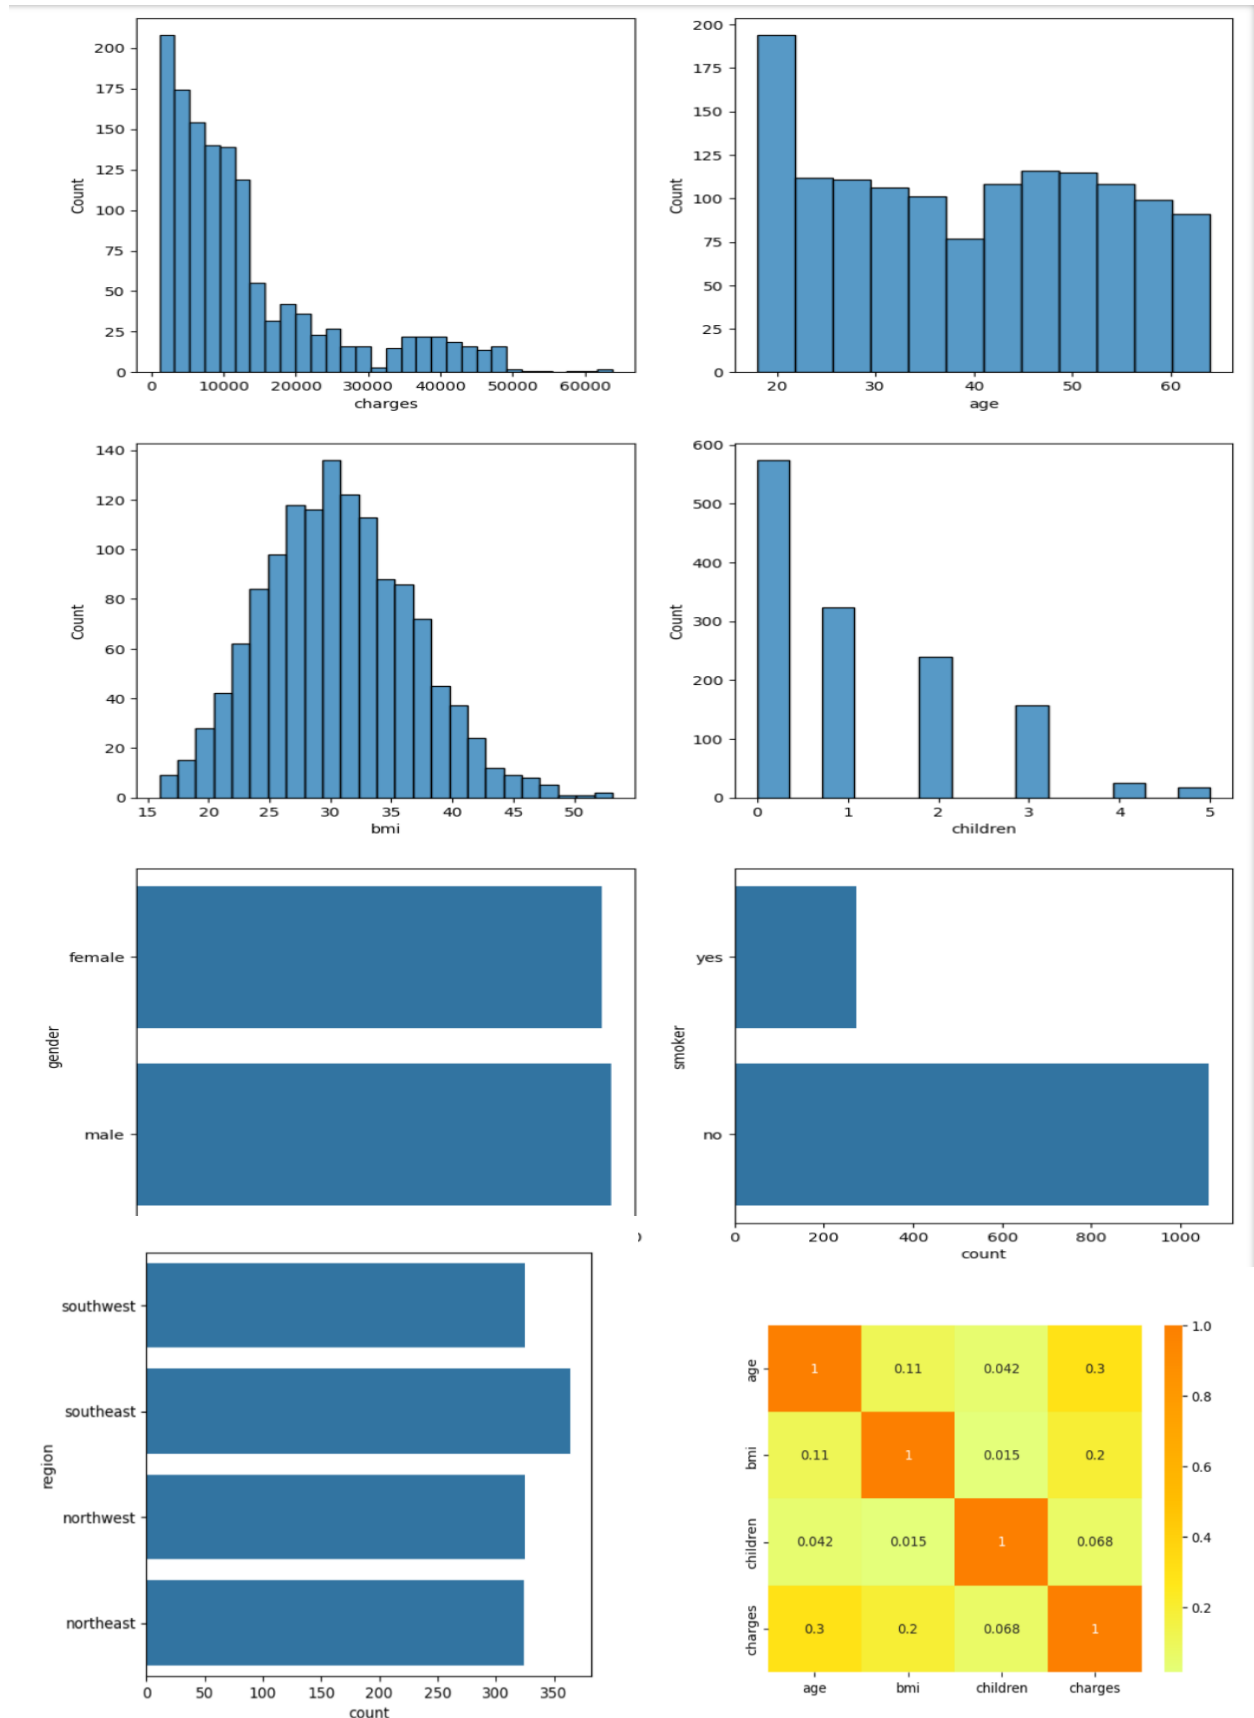
```
[62]:
```

```
#check how many values are missing (NaN) - after we filled in the NaN
count_nan = data.isnull().sum() # the number of missing values for every column
print(count_nan[count_nan > 0])
```

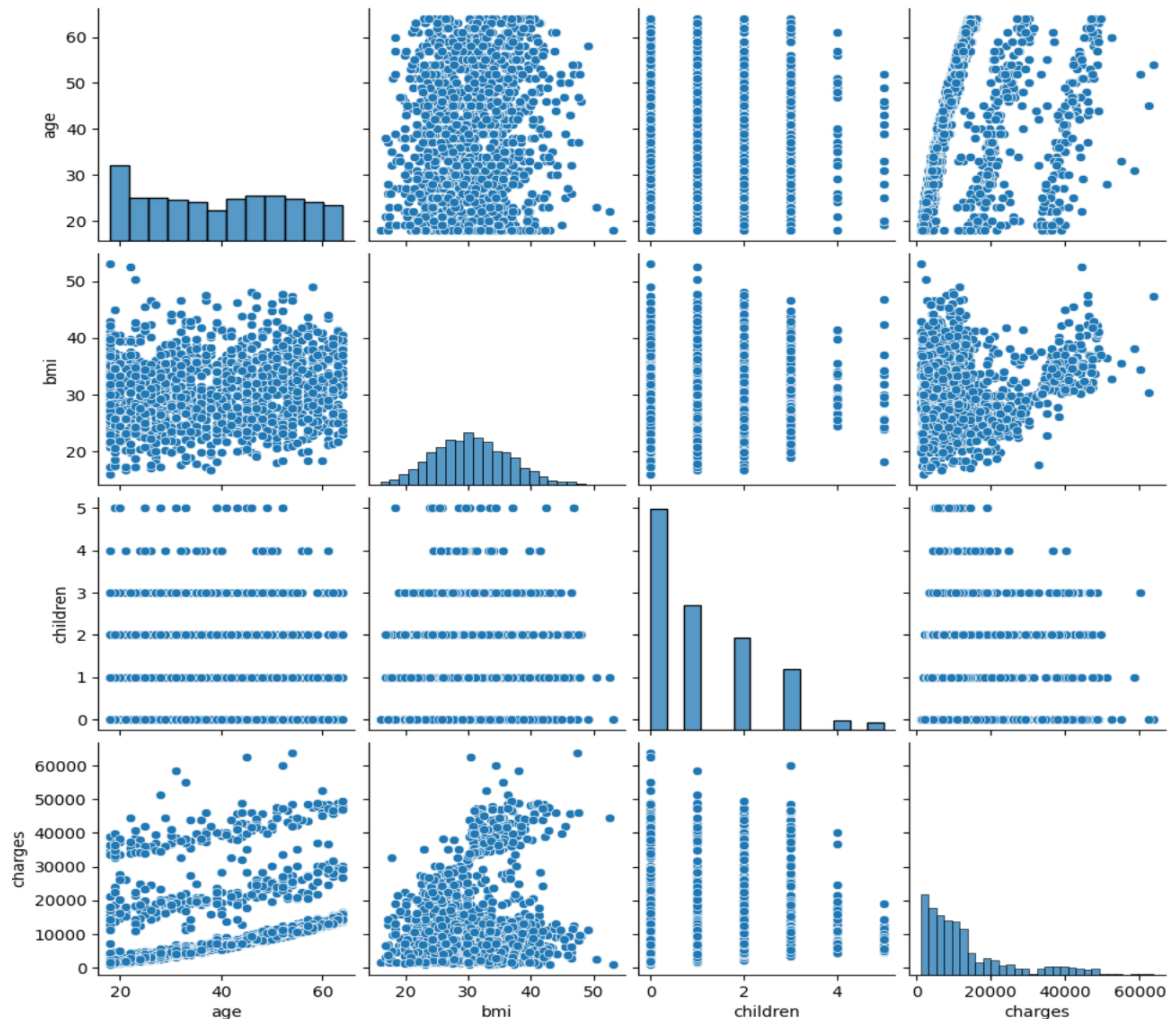
```
Series([], dtype: int64)
```

## Exploratory Data Analysis

First the report visualized independent variables by count to get some insights. As in following graphs, there are almost equal number of women and men were counted, almost half of people had no children. There are more non-smokers detected as well. At the following page, correlation has been plotted between numerical variables.

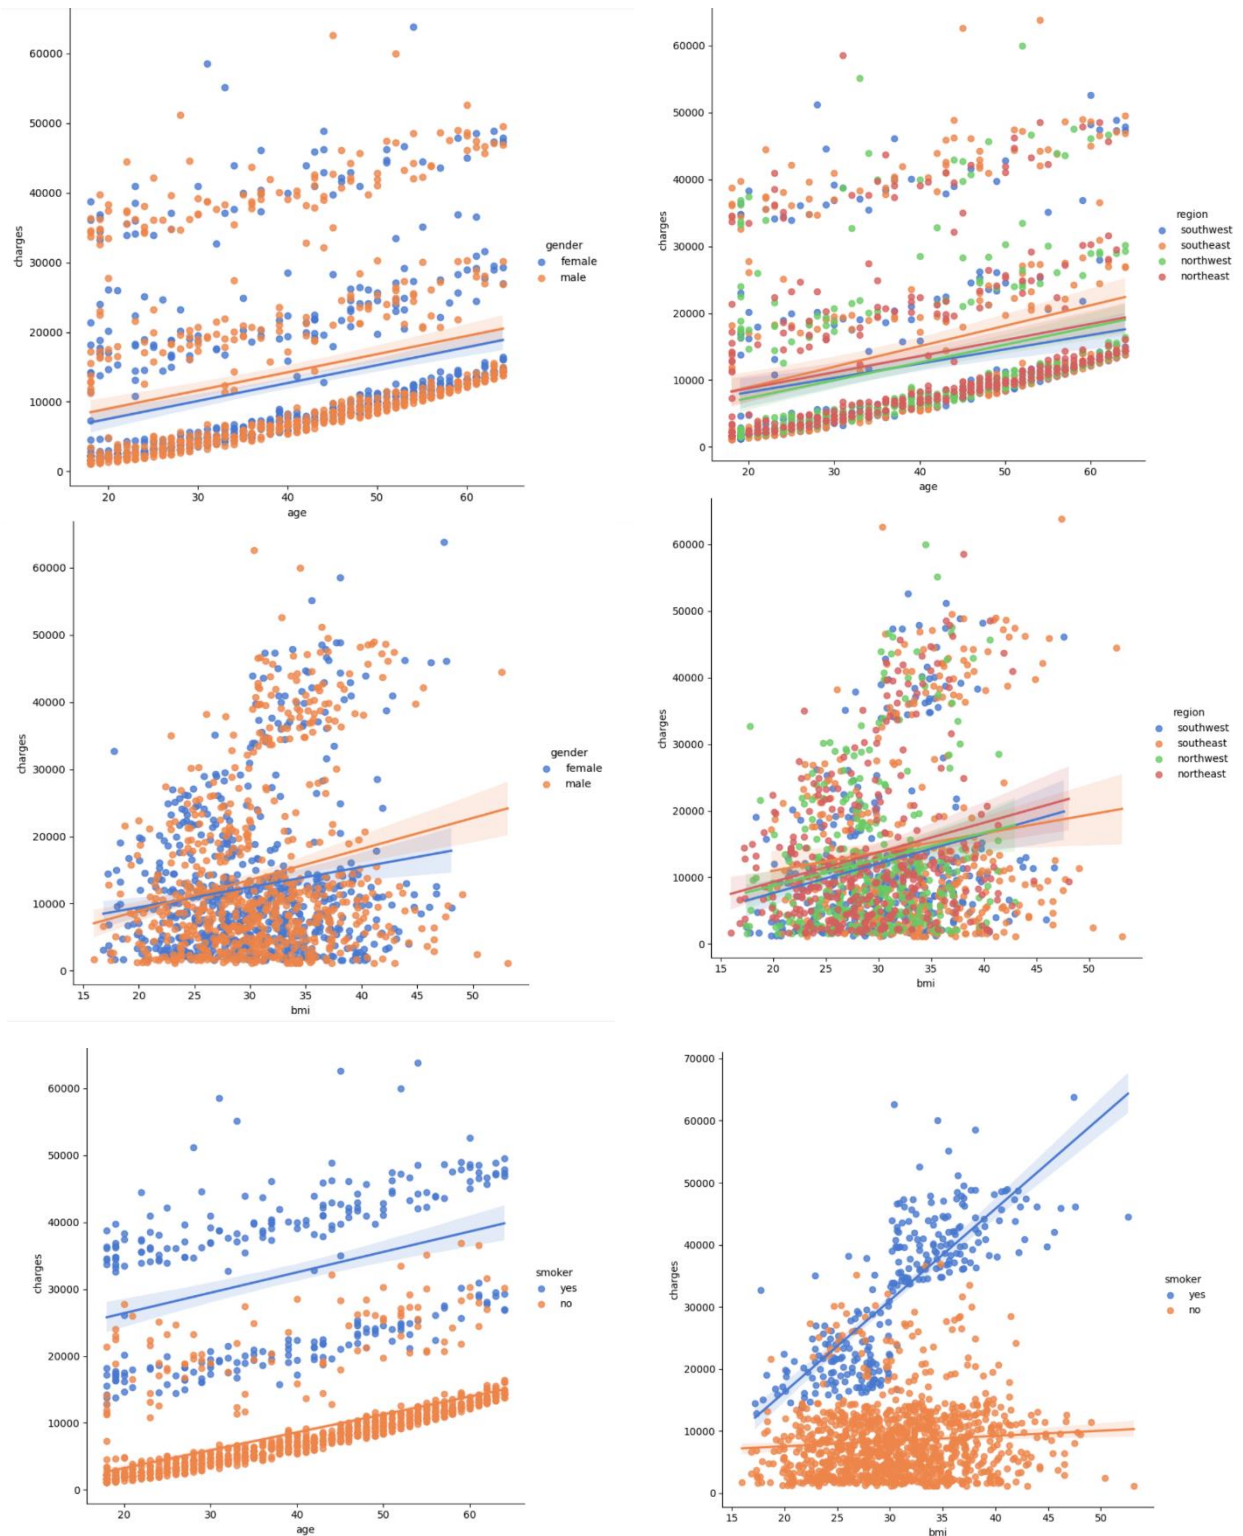


Next, the report further used pairplot between numeric variables in Python to see which numerical variables are affecting the charges more.



Looking at graphs above and correlation matrix, age vs charges and BMI vs charges has some patterns and higher correlations that can be further investigated by adding third categorical variables: smoker, region, and gender. Therefore, it added third categorical variables to see their effects on charges.

When region and gender were added to age vs charges and BMI vs charges plots, the report hasn't found clear linear distinction. However, smoker categorical variable was affecting these two plots and showed clear linear relationships which will improve our predictive model well. (Graphs are following)



Medical bills can be significantly higher regardless of patient's age for smokers and for smokers as BMI increases medical costs were increased. For non-smokers as BMI increases, its medical costs haven't



increased. Therefore, being a smoker can increase charges and having a higher BMI and smoking also can increase charges.

### Comparison of predictive models

The report ran five regression models and see which one performs better to predict the y variable.

Each model's training and testing followed these steps:

1. Splits the data into training and testing sets. Testing portion was set at 33 percent on each model, which prevents overfit and using separate parts of the data for training and testing helps to ensure the reliability and generalizability of your machine learning models.
2. Scales the features using Standard Scaler.
3. Fit the model
4. Makes predictions on the training and testing sets.
5. Prints the training and testing scores.
6. Evaluates the model using Mean Squared Error and R-squared.
7. Visualizes the predicted vs. true values for both the training and test sets.

In order to train and test data all categorical variables were turned into dummy variables since regression algorithms only accept numerical variables. Following encoding was done on Python:

```

1  #sklearn one hot encoding: maps each category to 0 (cold) or 1 (hot)
2
3  #one hot encoder = ohe
4  #create ndarray for one hot encoding (sklearn)
5  region = data.iloc[:,5:6].values #ndarray
6
7  ## ohe for region
8  ohe = OneHotEncoder()
9
10 region = ohe.fit_transform(region).toarray()
11 region = pd.DataFrame(region)
12 region.columns = ['northeast', 'northwest', 'southeast', 'southwest']
13 print("Sklearn one hot encoder results for region:")
14 print(region[:10])
15
16
17 #sklearn Label encoding: maps each category to a different integer
18
19 #create ndarray for label encoding (sklearn)
20 gender = data.iloc[:,1:2].values
21 smoker = data.iloc[:,4:5].values
22
23 #label encoder = le
24 ## le for gender
25 le = LabelEncoder()
26 gender[:,0] = le.fit_transform(gender[:,0])
27 gender = pd.DataFrame(gender)
28 gender.columns = ['gender']
29 le_gender_mapping = dict(zip(le.classes_, le.transform(le.classes_)))
30 print("Sklearn label encoder results for gender:")
31 print(le_gender_mapping)
32 print(gender[:10])
33
34 ## le for smoker
35 le = LabelEncoder()
36 smoker[:,0] = le.fit_transform(smoker[:,0])
37 smoker = pd.DataFrame(smoker)
38 smoker.columns = ['smoker']
39 le_smoker_mapping = dict(zip(le.classes_, le.transform(le.classes_)))
40 print("Sklearn label encoder results for smoker:")
41 print(le_smoker_mapping)
42 print(smoker[:10])

```



All coded variables and machine learning ready data are printed in Python as following:

```
[26]: #putting the data together:

##take the numerical data from the original data
X_num = data[['age', 'bmi', 'children']].copy()

##take the encoded data and add to numerical data
X_final = pd.concat([X_num, region, gender, smoker], axis = 1)

#define y as being the "charges column" from the original dataset
y_final = data[['charges']].copy()

#print all columns that will be used for testing
print(X_final,y_final)

#Test train split
X_train, X_test, y_train, y_test = train_test_split(X_final, y_final, test_size = 0.33, random_state = 0 )
#X_train, X_test, y_train, y_test = train_test_split(data[['age']], y_final, test_size = 0.33, random_state = 0 )
```

	age	bmi	children	northeast	northwest	southeast	southwest	\
0	19	27.900	0	0.0	0.0	0.0	1.0	
1	18	33.770	1	0.0	0.0	1.0	0.0	
2	28	33.000	3	0.0	0.0	1.0	0.0	
3	33	22.705	0	0.0	1.0	0.0	0.0	
4	32	28.880	0	0.0	1.0	0.0	0.0	
...	...	...	...	...	...	...	...	
1333	50	30.970	3	0.0	1.0	0.0	0.0	
1334	18	31.920	0	1.0	0.0	0.0	0.0	
1335	18	36.850	0	0.0	0.0	1.0	0.0	
1336	21	25.800	0	0.0	0.0	0.0	1.0	
1337	61	29.070	0	0.0	1.0	0.0	0.0	

	gender	smoker
0	0	1
1	1	0
2	1	0
3	1	0
4	1	0
...	...	...
1333	1	0
1334	0	0
1335	0	0
1336	0	0
1337	0	1



	charges
0	16884.92400
1	1725.55230
2	4449.46200
3	21984.47061
4	3866.85520
...	...
1333	10600.54830
1334	2205.98080
1335	1629.83350
1336	2007.94500
1337	29141.36030

[1338 rows x 9 columns]

[1338 rows x 1 columns]


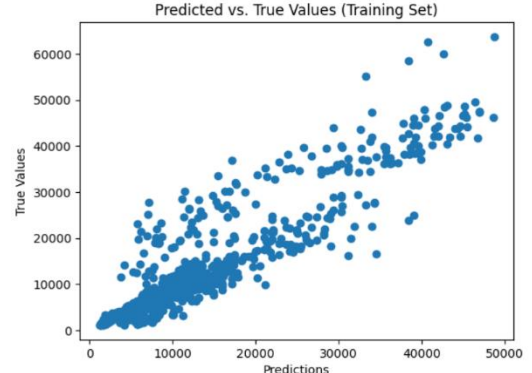
Next, the report trained and tested these models and evaluated on which one performs best to predict y variable. Each model has been visualized with train score and test score, mean squared error and r squared printed at the top of the plots.

Train score: 0.728, Test score: 0.786  
Mean Squared Error: 34338314.573869  
R-squared: 0.7855951871694039




Linear regression

Train score: 0.837, Test score: 0.802  
Mean Squared Error: 31634984.011629965  
R-squared: 0.8024744979453954





K nearest neighbors

Train score: 0.974, Test score: 0.859  
Mean Squared Error: 22542924.43349883  
R-squared: 0.8592443585598498




Random forest regression

poly train score 0.828, poly test score: 0.870  
Mean Squared Error: 24087593.802425988  
R-squared: 0.79181376159658



Polynomial regression

Train score: 0.999, Test score: 0.701  
 Mean Squared Error: 47808139.881123416  
 R-squared: 0.7014910192828288



Decision Tree

	Model	Train Score	Test Score	Mean Squared Error	\
0	Linear Regression	0.728	0.786	3.433831e+07	
1	KNN	0.837	0.802	3.163498e+07	
2	Random Forest	0.974	0.859	2.254292e+07	
3	Polynomial Regression	0.828	0.870	2.408759e+07	
4	Decision Tree	0.999	0.701	4.780814e+07	


  

	R-squared
0	0.7856
1	0.8024
2	0.8592
3	0.7918
4	0.7014

At the chart above, it summarized all scores in comparison between models. It looks like random forest regression best predicts y variable at 86 percent with less error. Another good predictor can be polynomial regression model but it has higher errors, lower r squared, and lower training score.

Feature ranking:

1. Feature 8 (0.623069)
2. Feature 1 (0.205051)
3. Feature 0 (0.130234)
4. Feature 2 (0.018765)
5. Feature 7 (0.005492)
6. Feature 4 (0.005227)
7. Feature 3 (0.005140)
8. Feature 5 (0.003740)
9. Feature 6 (0.003281)



On random forest model, the report visualized features importance and ranked it. Since Python starts counting from zero, #8 is the smoker feature, #1 is the BMI feature, #0 is the age, #2 is number of dependents feature (Page 8) which affected y variable prediction higher than other variables such as region and gender.

### **Executive summary**

The aim of this analysis was to investigate the predictability of medical costs (denoted as 'y' variable) based on various factors including age, BMI, gender, geographical regions, number of dependents, and smoking habits. The study explores how accurately these factors can forecast medical expenses, which could assist health insurance companies in effectively predicting costs and tailoring premium fees accordingly, thus optimizing cost efficiency.

By employing predictive modeling techniques, particularly Random Forest Regression, it was determined that the model achieved a prediction accuracy of 86 percent. This indicates a high level of predictability for medical costs based on the specified input variables. Therefore, health insurance companies can utilize these findings to enhance their risk assessment strategies, better manage financial resources, and offer more customized premium fees to their customers.

In conclusion, the insights derived from this analysis provide valuable guidance for health insurance companies to make informed decisions, optimize their operations, and ultimately improve the overall efficiency of their business processes.

## **Appendix**

1. Data: <https://www.kaggle.com/datasets/mirichoi0218/insurance/data>

I hasn't followed the codes thoroughly, but took some ideas from the codes listed and ran Python codes to best illustrate the data and best illustrate prediction models.



Final project in  
Python.html

2. Python codes: